

# Evaluation of segmentation methods based on the BSDS500 benchmark

Edgar A. Margffoy-Tuay  
Universidad de los Andes  
201412566

ea.margffoy10@uniandes.edu.co

## Abstract

*To evaluate a set of algorithms designed to solve a specific problem of interest, it is essential to define a common ground of evaluation, such as defining a dataset proposed to evaluate the problem, this database shall contain a set of inputs, labels and a evaluation methodology that allows different teams to compare their solutions to the problem to other proposals. In the present document, we pretend to introduce the Berkeley Segmentation Dataset and Benchmarks (BSDS500) as a framework to evaluate the performance of different segmentation algorithms, and more specifically, we compare the performance of different unsupervised segmentation algorithms, like *k*-means, watershed and UCM, by comparing their results on the Precision-Recall curve defined on BSDS500.*

## 1. Introduction

Before the introduction of normalized benchmarks to evaluate the performance of different algorithms proposed to solve an specific task, each research team chose a set of predefined catalog images, on which, the proposed method results were evaluated and presented, this evaluation pipeline represented a challenge when different segmentation algorithms were to be compared and tested against, which prompted an uneven challenge between the different teams. To account

for this difficulties, it was essential to define a benchmark that provides a set of uniform and public images with their respective labels and annotations, along with the dataset, an evaluation methodology must be provided, this framework allows to compare and classify different algorithms on a common ground.

On image segmentation context, the preferred benchmark was proposed by the Computer Vision Group at UC Berkeley, their database (BSDS500) [1], consisted of 500 images (Photographs), which were labeled by different humans. All the segmentations defined on the database have different boundary and region precision, *i.e.*, There are more or less defined regions per mask annotation. This allows to evaluate the segmentation using the Precision-Recall curve and the  $F_1$  score statistic, which in turn can be compared with the Human score evaluated over the image segmentation task.

## 2. Materials and Methods

On the present document, a comparison of different unsupervised segmentation methods is proposed, based on the BSDS500 benchmark evaluation pipeline, the methods to be compared against correspond to *k*-means and Ultrametric Contour Maps, each model parameters are setup based on the results obtained after the heuristic evaluation done over a subset of 20 images of the dataset, the evaluation score was different from

the evaluation methodology proposed by the UC Berkeley team (Precision/Recall) and was based on contour comparison of the results against the ground truth contour masks. Now the results are evaluated over all the test set of BSDS500 dataset, using the fast morphological benchmark, which compares the boundary reconstruction of each image by using the segmentation morphological reconstruction of the ground truth boundary masks. To represent the input images, each input vector was represented as a coordinate on Lab space, with spatial information. The initial parameters were evaluated over the validation and training sets of the database, the chosen number of segmentation regions correspond to 5, 10, 20, 50 and 100, mapping to 5 points on the AP curve respectively. With respect to the watershed

The implementation of the watershed procedure is based on the OpenCV<sup>1</sup> Python bindings, with the regional minima extraction routines adapted from the Gala library<sup>2</sup>, the clustering methods are based on the Scikit-learn<sup>3</sup> implementations. Also, the centroid and means initialization for K-Means was based on the KMC<sup>24</sup> algorithm [2], all the routines were implemented on Cython, to improve execution times and the efficiency of the evaluation procedure.

## About the models

### Watershed Transform

The watershed transform is based on the concept of watershed lines, these lines correspond to the local basin in which all drops of water are gathered when they fall down from a peak. In this context, the peaks of the images correspond to contours, and the space between borders may correspond to a river basin, the idea of the algorithm is to flood the image from the river basins (Regional minima), up to the peaks, at the end of the procedure, only the peaks remain, which implies that a boundary segmenta-

tion of the image is achieved. However, due to the existence of multiple regional minima on the gradient, this method is prone to produce an oversegmentation of the image, to overcome this limitation, it is possible to define a set of precomputed markers and impose them over the image, by executing this procedure, the flooding is not done over the set of markers, reducing the oversegmentation on the image. Due to the inexistence of precomputed markers over each image, the markers are selected as the regional minima of height  $h$  (h-minima), the value of  $h$  is increased, starting on 0, until the total number of regions is equal to  $K$ .

### K-Means

In the present experiment, K-Means was employed to segment an image by conforming clusters from pixel information, the clusters means were calculated by using the  $\ell_2$  distance metric (Euclidean), which gives rise to spherical and convex color clusters, which can be spatially sparse or localized, depending on the inclusion of spatial information. To segment an input image, for each color triplet, the label of the closest mean is assigned, finally, the segmented image is reshaped into the image original dimensions.

### Ultrametric Contour Maps

The Ultrametric Contour Maps is a hierarchical segmentation representation based on the Ultrametric Distance, it groups a set of different segmentations, based on the region contours and a similarity function between adjacent regions. This similarity distance must comply with the ultrametric inequality, for instance, a distance between points on Lab\* space can be define as a mesure of dissimilarity of regions on the hierarchy, this definition allows to refine the contour detection and segmentation of an image. Also, UCM compiles all the segmentation trees on a single contour image, which can be thresholded to extract each of the individual segmentation levels present on the ultrametric dendrogram.

<sup>1</sup><http://opencv.org>

<sup>2</sup><https://github.com/janelia-flyem/gala>

<sup>3</sup><http://scikit-learn.org/stable/>

<sup>4</sup><https://github.com/obachem/kmc2>

### 3. Results

Before the evaluation of all segmentation methods, it was expected that UCM outperforms the performance of both k-means and watersheds, due to the expressivity implied by using a segmentation hierarchy family instead of a single segmentation, this approach guarantees that more information is captured on a single image, compared to the clustering based approaches, such as K-Means or morphological methods like watersheds, that do not take in account more relations among pixels and regions and are very limited. Among the strengths of UCM we can found, it is that the ultrametric distance between regions impose more restrictions and conditions that the regions must comply in order to be merged, those restrictions expressed in terms of color, border and distance affinity can be regularized and controlled via hyperparameters, which can fine-tune the results of the method.

On the other hand, the hyperparameters defined on the clustering methods, *i.e.*,  $K$ , just only define the number of segmentation regions to be found, however, this parameter does not define or expresses how these regions should be found. This limitation accounts to the fact that unsupervised segmentation methods, such as k-means are based on the fact that grouping should be defined and built upon a distance that relates the elements on the image space, nevertheless, there does not exists a deterministic schema to define the representation of each point and the distance definition between each point on that space. In conclusion, is not possible to regularize nor guide the segmentation process if those methods are to be used.

As it can be seen on Figures 1 and 2, the overall PR<sup>5</sup> curve comparison between the previously defined methods allows to draw the same conclusion as the exposed previously, in this case, the UCM performance is far from better than the

<sup>5</sup>In the present benchmark, it is desired that the  $F_1$  score approximates or overshoots the human performance (0.79) and approximate to the value 1, on which all instances are segmented correctly (Precision) and no correct segmentation is missed (Recall)

results obtained after evaluating the PR scores for k-means, as it can be inferred of the benchmark evaluation statistics (Table 1), by which, the k-means  $F_1$  score (0.54) is approximately lesser on 20% than the score obtained after evaluating UCM (0.73). This implies that UCM outperforms k-means, due to the expressive and rich hierarchical representation, which adapt to the different segmentation levels and thresholds and segmentation families according to the ultrametric distance and the dissimilarity measure between regions, which increases both Precision and Recall, compared to K-means, which present an average Recall but a low precision, which means that none of the segmentation boundaries are lost, but the segmentation proposed does not right. It is worth to notice that evaluating Watershed performance was difficult due to the uniform evaluation of the number of regions chosen on each image, *i.e.*, Some images could present 1000 regions, but others only 999 regions, effect that renders the comparison useless and biased against a subset of the values subject to evaluation, and therefore, the results were not subject to comparison.

| Method     | ODS  | OIS  | AP   |
|------------|------|------|------|
| <b>UCM</b> | 0.73 | 0.76 | 0.72 |
| K-Means    | 0.54 | 0.57 | 0.3  |

Table 1: Precision - Recall evaluation results on BSDS500 of UCM and K-means over the test set




Figure 1: Precision-Recall curve results associated to UCM evaluation over BSDS500 test set




Figure 2: Precision-Recall curve results associated to K-means (Lab+xy) evaluation over BSDS500 test set

## 4. Conclusions

After evaluating different unsupervised segmentation methods, it is possible to appreciate that

those that have more regularization parameters, present more refined segmentation regions and boundary contours, due to the definition of higher-level distance metrics and similarity measures between groups of pixels, increasing the number of parameters and affinities allows us to improve the overall result when a method is evaluating on a common ground benchmark like BSDS500. The methods based only on clustering are limited in terms of regularization and control of the segmentation regions to define, *i.e.*, It is not possible to guide the segmentation regions that shall result after applying one of those methods, also, the computational time complexity increases if the number of desired regions increases, which only decreases the overall accuracy score on the PR curve. Finally, the models that exploit the expressiveness of a hierarchical segmentation approach can capture more information, and therefore present a better performance when subject to evaluation on different benchmarks, like BSDS500. To improve the accuracy of some of the clustering methods, it should be possible to enrich and extend the representation space of each pixel on the image and the distance definition between groups as well, however this shall increase the time complexity and may not approximate to better and expressive models, such as UCM.

## References

- [1] P. Arbelaez, M. Maire, C. Fowlkes, and J. Malik. Contour detection and hierarchical image segmentation. *IEEE Trans. Pattern Anal. Mach. Intell.*, 33(5):898–916, May 2011.
- [2] O. Bachem, M. Lucic, S. H. Hassani, and A. Krause. Approximate k-means++ in sublinear time. In *Proceedings of the Thirtieth AAAI Conference on Artificial Intelligence*, AAAI’16, pages 1459–1467. AAAI Press, 2016.